

Numpy

Numerical Computing in Python

What is Numpy?

- Numpy, Scipy, and Matplotlib provide MATLAB-like

functionality in python.

- Numpy Features:
 - Typed multidimensional arrays (matrices)
 - Fast numerical computations (matrix math)
 - High-level math functions

Why do we need NumPy

Let's see for ourselves!

Why do we need NumPy

- Python does numerical computations slowly.
- 1000 x 1000 matrix multiply
 - Python triple loop takes > 10 min.
 - Numpy takes ~0.03 seconds

NumPy Overview

1. Arrays
2. Shaping and transposition
3. Mathematical Operations
4. Indexing and slicing
5. Broadcasting

Arrays

Structured lists of numbers.

- Vectors
- Matrices
- Images
- Tensors
- ConvNets

Arrays


Structured lists of
numbers. • **Vectors**

• **Matrices**

• Images

• Tensors

• ConvNets



Arrays

Structured lists of numbers.



- Vectors
- Matrices
- **Images**
- Tensors
- ConvNets



Arrays

Structured lists of numbers.

- Vectors
- Matrices
- Images
- **Tensors**
- **ConvNets**



Arrays

Structured lists of numbers.

- Vectors
- Matrices
- Images
- Tensors
- ConvNets

MATRICES



IMAGES



TENSORS



CONVNETS



Arrays, Basic Properties

```
import numpy as np
a = np.array([[1, 2, 3], [4, 5, 6]], dtype=np.float32) print
(a.ndim, a.shape, a.dtype )
```

1. Arrays can have any number of dimensions, including zero (a scalar).
2. Arrays are typed: np.uint8, np.int64, np.float32, np.float64
3. Arrays are dense. Each element of the array exists and has the same type. ¹¹

Arrays, creation

- `np.ones`, `np.zeros`
- `np.arange`
- `np.concatenate`
- `np.astype`
- `np.zeros_like`, `np.ones_like`
- `np.random.random`

Arrays, creation

- **np.ones, np.zeros**
- np.arange
- np.concatenate
- np.astype
- np.zeros_like,
np.ones_like
- np.random.random

```
>>> np.ones((3,5),dtype=np.float32)
array([[ 1.,  1.,  1.,  1.,  1.],
       [ 1.,  1.,  1.,  1.,  1.],
       [ 1.,  1.,  1.,  1.,  1.]], dtype=float32)
```

```
>>> np.zeros((6,2),dtype=np.int8)
array([[0, 0],
       [0, 0],
       [0, 0],
       [0, 0],
       [0, 0],
       [0, 0]], dtype=int8)
```

Arrays, creation

- np.ones, np.zeros
- **np.arange**
- np.concatenate
- np.astype
- np.zeros_like, np.ones_like
- np.random.random

```
>>> np.arange(1334,1338)  
array([1334, 1335, 1336, 1337])
```

Arrays, creation

- np.ones, np.zeros
- np.arange
- **np.concatenate**
- np.astype
- np.zeros_like,
np.ones_like

```
>>> A = np.ones((2,3))
>>> B = np.zeros((4,3))
>>> np.concatenate([A,B])
array([[ 1.,  1.,  1.],
       [ 1.,  1.,  1.],
       [ 0.,  0.,  0.],
       [ 0.,  0.,  0.],
       [ 0.,  0.,  0.],
       [ 0.,  0.,  0.]])
```

- np.random.random

15

Arrays, creation

- np.ones, np.zeros
- np.arange
- **np.concatenate**
- np.astype
- np.zeros_like,

```
>>> A = np.ones((4,1))
>>> B = np.zeros((4,2))
>>> np.concatenate([A,B], axis=1)
array([[ 1.,  0.,  0.],
       [ 1.,  0.,  0.],
       [ 1.,  0.,  0.],
       [ 1.,  0.,  0.]])
```


np.ones_like

- np.random.random

16

Arrays, creation

- np.ones, np.zeros
- np.arange
- np.concatenate
- **np.astype**

```
>>> A
array([[ 4670.5,  4670.5,  4670.5],
       [ 4670.5,  4670.5,  4670.5],
       [ 4670.5,  4670.5,  4670.5],
       [ 4670.5,  4670.5,  4670.5],
       [ 4670.5,  4670.5,  4670.5]], dtype=float32)
>>> print(A.astype(np.uint16))
[[4670 4670 4670]
 [4670 4670 4670]
 [4670 4670 4670]
```

- np.zeros_like, np.ones_like
- np.random.random

17

Arrays, creation

- np.ones, np.zeros
- np.arange
- np.concatenate
- np.astype

```
>>> a = np.ones((2,2,3))  
>>> b = np.zeros_like(a)  
>>> print(b.shape)
```

- **np.zeros_like, np.ones_like**
- np.random.random

18

Arrays, creation

- np.ones, np.zeros
- np.arange
- np.concatenate
- np.astype

```
>>> np.random.random((10,3))
array([[ 0.61481644,  0.55453657,  0.04320502],
       [ 0.08973085,  0.25959573,  0.27566721],
       [ 0.84375899,  0.2949532 ,  0.29712833],
       [ 0.44564992,  0.37728361,  0.29471536],
       [ 0.71256698,  0.53193976,  0.63061914],
       [ 0.03738061,  0.96497761,  0.01481647],
       [ 0.09924332,  0.73128868,  0.22521644],
       [ 0.94249399,  0.72355378,  0.94034095],
```

- `np.zeros_like`, `np.ones_like`
- **`np.random.random`**

19

Arrays, danger zone

- Must be dense, no holes.
- Must be one type
- Cannot combine arrays of different shape

```
>>> np.ones([7,8]) + np.ones([9,3])
Traceback (most recent call last):
  File "<stdin>", line 1, in <module>
ValueError: operands could not be broadcast together
with shapes (7,8) (9,3)
```

20

Shaping

```
a = np.array([1, 2, 3, 4, 5, 6])
```

```
a = a.reshape(3, 2)
```

```
a = a.reshape(2, -1)
```

```
a = a.ravel()
```

1. Total number of elements cannot change.
2. Use -1 to infer axis shape
3. Row-major by default (MATLAB is column-major)

21

Transposition

```
a = np.arange(10).reshape(5, 2)
```

```
a = a.T
```

```
a = a.transpose((1, 0))
```

`np.transpose` permutes axes.

`a.T` transposes the first two axes.

22

Saving and loading arrays

```
np.savez('data.npz', a=a)
```

```
data = np.load('data.npz')
```

```
a = data[ 'a' ]
```

1. NPZ files can hold multiple arrays
2. np.savez_compressed similar.

Image arrays

Images are 3D arrays: width, height, and channels
Common image formats:



height x width x RGB (band-interleaved)

height x width (band-sequential)

Gotchas:

Channels may also be BGR (OpenCV does this)

May be [width x height], not [height x width]

24

Saving and Loading Images

SciPy: `skimage.io.imread`, `skimage.io.imsave`

height x width x RGB

PIL / Pillow: `PIL.Image.open`, `Image.save`

width x height x RGB

OpenCV: `cv2.imread`, `cv2.imwrite`

height x width x BGR

Why Pillow?

The [Python Imaging Library](#), (PIL) is the library
for image manipulation, however...

Why Pillow?

... PIL's last release was in

2009



Why Pillow?

- easier to install ●
- supports Python 3 ●
- active development ●
- actually works*

What can Pillow do for me?

- Automatically generate thumbnails
- Apply image filters (auto-enhance)
- Apply watermarks (alpha layers) ●
- Extract images from animated gifs ●
- Extract image metadata
- Draw text for annotations (and shapes) ●
- Basically script things that you might do in Photoshop or GIMP for large numbers of images, in

Python

- ImageOps
- ImageMath • ImageFilter • ImageEnhance • ImageStat

Modules:

Pillow Setup

Pillow's prerequisites:

<https://pypi.python.org/pypi/Pillow/2.1.0#platform-specific-instructions>

Warning!

Since some (most?) of Pillow's features require external libraries, prerequisites

can be a little tricky

After installing prereqs:

```
$ pip install Pillow
```

Documentation

Pillow documentation:

<http://pillow.readthedocs.org/en/latest/about.html>

Image Basics:

<http://pillow.readthedocs.org/en/latest/handbook/concepts.html>

List of Modules:

<http://pillow.readthedocs.org/en/latest/reference/index.html>

Read in an image!

```
from PIL import Image

im = Image.open(infile)
im.show()

print(infile, im.format, "%dx%d" %
im.size, im.mode)
```


Basic Methods

geometric

transforms:

```
out = im.resize((128, 128), Image.ROTATE_45, Image.ResampleMethod.BILINEAR)
out = im.rotate(45, Image.ROTATE_45, Image.ResampleMethod.BILINEAR)
im.resize((512, 512), Image.ROTATE_45, Image.ResampleMethod.BILINEAR)
```

```
out = im.transpose(Image.FLIP_LEFT_RIGHT)
out = im.transpose(Image.FLIP_TOP_BOTTOM)
```

crop:

```
box = (100, 100, 400, 400) #(left, upper, right, lower)  
region = im.crop(box)
```

Recap

We just saw how to create arrays, reshape them, and

permute axes Questions so far?

Recap

We just saw how to create arrays, reshape them, and

permute axes Questions so far?

Now: let's do some math

35

Mathematical operators

- Arithmetic operations are element-wise
- Logical operator return a bool array


- In place operations modify the array

36

Mathematical operators

- **Arithmetic operations are element-wise**
- Logical operator return a bool array

- In place operations modify the array



37

Mathematical operators

- Arithmetic operations are element-wise
- **Logical operator return a bool array**

- In place operations modify the array




38

Mathematical operators

- Arithmetic operations are element-wise

- Logical operator return a bool array
- **In place operations modify the array**



Math, upcasting

Just as in Python and Java, the result of a math operator is cast to the more general or precise datatype.

`uint64 + uint16 => uint64`

`float32 / int32 => float32`

Warning: upcasting does not prevent overflow/underflow. You must manually cast first.

Use case: images often stored as `uint8`. You should convert to `float32` or `float64` before doing math.

Math, universal functions

Also called ufuncs

Element-wise

Examples:

- `np.exp`
- `np.sqrt`
- `np.sin`
- `np.cos`
- `np.isnan`

Math, universal functions

Also called ufuncs

Element-wise

Examples:

- `np.exp`
- `np.sqrt`
- `np.sin`
- `np.cos`



- np.isnan

42

Math, universal functions

Also called ufuncs

Element-wise

Examples:

- np.exp
- np.sqrt
- np.sin



- np.cos
- np.isnan

43

Indexing

`x[0,0]` # top-left element

`x[0,-1]` # first row, last column

`x[0,:]` # first row (many entries)

`x[:,0]` # first column (many entries)

Notes:

- Zero-indexing
- Multi-dimensional indices are comma-separated (i.e., a tuple) 44

Indexing, slices and arrays

```
I[1:-1,1:-1] # select all but one-pixel  
border
```

```
I = I[:, :, ::-1] # swap channel order
```

```
I[I<10] = 0 # set dark pixels to black
```

```
I[[1,3], :] # select 2nd and 4th row
```

1. Slices are **views**. Writing to a slice overwrites the original array.
2. Can also index by a list or boolean array.

45

Python Slicing

Syntax: start:stop:step

```
a = list(range(10))
```

```
a[:3] # indices 0, 1, 2
```

```
a[-3:] # indices 7, 8, 9
```

```
a[3:8:2] # indices 3, 5, 7
```

```
a[4:1:-1] # indices 4, 3, 2 (this one is tricky) 46
```

Axes

```
a.sum() # sum all entries
```

```
a.sum(axis=0) # sum over rows
```

```
a.sum(axis=1) # sum over columns
```

```
a.sum(axis=1, keepdims=True)
```


1. Use the axis parameter to control which axis NumPy operates on
2. Typically, the axis specified will disappear, keepdims keeps all dimensions

47

Broadcasting

```
a = a + 1 # add one to every element
```

When operating on multiple arrays, broadcasting rules are used. Each dimension must match, from right-to-left

1. Dimensions of size 1 will broadcast (as if the value was repeated). 2. Otherwise, the dimension must have the same shape. 3. Extra dimensions of size 1 are added to the left as needed.

48

Broadcasting example

Suppose we want to add a color value to an image

a.shape is 100, 200, 3

b.shape is 3

a + b will pad b with two extra dimensions so it has an effective shape of 1 x 1 x 3.

So, the addition will broadcast over the first and second

dimensions. 49

Broadcasting failures

If `a.shape` is 100, 200, 3 but `b.shape` is 4 then `a + b` will fail. The trailing dimensions must have the same shape (or be 1)

Tips to avoid bugs

1. Know what your datatypes are.

2. Use matplotlib for sanity checks.
3. Know np.dot vs np.mult.